

# Class 7: Machine Learning 1

Ani A16647613


## Clustering Methods

The broad goal here is to find groupings (clusters) in your input data.

##K means

First, lets make up some data to cluster.

```
x <- rnorm(1000)
hist(x)
```



Make a vector of length 60 with 30 points centered at -3 and 30 points at +3.

```
tmp <- c(rnorm(30, -3), rnorm(30,3))
tmp
```

```
[1] -4.627177 -3.026079 -2.597380 -1.932038 -2.037396 -4.159953 -2.144216
[8] -3.773737 -1.611692 -4.031699 -2.356316 -2.769998 -3.023750 -4.282115
[15] -1.071738 -2.885831 -2.852391 -2.189565 -2.204870 -4.049411 -4.353746
[22] -2.583500 -4.045128 -4.028440 -3.786038 -4.574942 -3.688313 -3.259243
[29] -3.538056 -2.257877  3.050598  2.324934  3.129494  2.765599  3.083536
[36]  1.804232  3.359325  2.827801  2.318489  3.525650  1.802170  1.294771
[43]  3.031259  3.232475  2.411985  3.144426  3.410688  2.825867  3.074261
[50]  1.356501  2.698256  2.165933  2.787521  3.349106  2.316083  2.286532
[57]  4.969153  2.409445  3.580665  1.893275
```


I will now make a smaller x and y dataset with 2 groups of points.

```
x <- cbind(x=tmp, y=rev(tmp))
x
```

```
      x      y
[1,] -4.627177  1.893275
[2,] -3.026079  3.580665
[3,] -2.597380  2.409445
[4,] -1.932038  4.969153
[5,] -2.037396  2.286532
[6,] -4.159953  2.316083
[7,] -2.144216  3.349106
[8,] -3.773737  2.787521
[9,] -1.611692  2.165933
[10,] -4.031699  2.698256
[11,] -2.356316  1.356501
[12,] -2.769998  3.074261
[13,] -3.023750  2.825867
[14,] -4.282115  3.410688
[15,] -1.071738  3.144426
[16,] -2.885831  2.411985
[17,] -2.852391  3.232475
[18,] -2.189565  3.031259
[19,] -2.204870  1.294771
[20,] -4.049411  1.802170
[21,] -4.353746  3.525650
[22,] -2.583500  2.318489
```

```
[23,] -4.045128  2.827801
[24,] -4.028440  3.359325
[25,] -3.786038  1.804232
[26,] -4.574942  3.083536
[27,] -3.688313  2.765599
[28,] -3.259243  3.129494
[29,] -3.538056  2.324934
[30,] -2.257877  3.050598
[31,]  3.050598 -2.257877
[32,]  2.324934 -3.538056
[33,]  3.129494 -3.259243
[34,]  2.765599 -3.688313
[35,]  3.083536 -4.574942
[36,]  1.804232 -3.786038
[37,]  3.359325 -4.028440
[38,]  2.827801 -4.045128
[39,]  2.318489 -2.583500
[40,]  3.525650 -4.353746
[41,]  1.802170 -4.049411
[42,]  1.294771 -2.204870
[43,]  3.031259 -2.189565
[44,]  3.232475 -2.852391
[45,]  2.411985 -2.885831
[46,]  3.144426 -1.071738
[47,]  3.410688 -4.282115
[48,]  2.825867 -3.023750
[49,]  3.074261 -2.769998
[50,]  1.356501 -2.356316
[51,]  2.698256 -4.031699
[52,]  2.165933 -1.611692
[53,]  2.787521 -3.773737
[54,]  3.349106 -2.144216
[55,]  2.316083 -4.159953
[56,]  2.286532 -2.037396
[57,]  4.969153 -1.932038
[58,]  2.409445 -2.597380
[59,]  3.580665 -3.026079
[60,]  1.893275 -4.627177
```

```
plot(x)
```



```
k <- kmeans(x, centers=2)
k
```

K-means clustering with 2 clusters of sizes 30, 30

Cluster means:

	x	y
1	2.741001	-3.124754
2	-3.124754	2.741001

Clustering vector:

```
[1] 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 1 1 1 1 1 1 1 1
[39] 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1
```

Within cluster sum of squares by cluster:

```
[1] 43.21223 43.21223
(between_SS / total_SS = 92.3 %)
```

Available components:

[1] "cluster"	"centers"	"totss"	"withinss"	"tot.withinss"
[6] "betweenss"	"size"	"iter"	"ifault"	

Q. From your result object `k` how many points are in each cluster?

```
k$size
```

```
[1] 30 30
```

Q. What “component” of your result object details the cluster membership?

```
k$cluster
```

```
[1] 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 1 1 1 1 1 1 1 1
[39] 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1
```


Q. Cluster centers?

```
k$centers
```

```
      x      y
1  2.741001 -3.124754
2 -3.124754  2.741001
```

Q. Plot of our clustering results

```
plot(x, col= k$cluster)
points(k$centers, col="blue", pch=15, cex=2)
```



Now to make it into 4 groups.

```
k4 <- kmeans(x, centers=4)
k4
```

K-means clustering with 4 clusters of sizes 30, 10, 7, 13

Cluster means:

	x	y
1	-3.124754	2.741001
2	3.338730	-2.452689
3	2.034808	-2.325283
4	2.661467	-4.072212

Clustering vector:

[1]	1	1	1	1	1	1	1	1	1	1	1	1	1	1	1	1	1	1	1	1	1	1	1	1	1	1	1	1	1	
[39]	3	4	4	3	2	2	3	2	4	2	2	3	4	3	4	2	4	3	2	3	2	4								

Within cluster sum of squares by cluster:

```
[1] 43.212230  7.266583  2.513883  5.625190
(between_SS / total_SS = 94.8 %)
```


```
[1] "cluster"      "centers"      "totss"        "withinss"     "tot.withinss"
[6] "betweenss"    "size"         "iter"         "ifault"
```

[1] 30 10 7 13

```
[1] 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 2 4 2 4 4 4 4 4
[39] 3 4 4 3 2 2 3 2 4 2 2 3 4 3 4 2 4 3 2 3 2 4
```

	x	y
1	-3.124754	2.741001
2	3.338730	-2.452689
3	2.034808	-2.325283
4	2.661467	-4.072212

7



A big limitation of `kmeans` is that it does what you ask even if you ask for silly clusters.

## Hierarchical Clustering

The main base R function for Hierarchical Clustering is `hclust()`. Unlike `kmeans()` you cannot just pass in your data as an input. You first need to calculate a distance matrix.

```
d <- dist(x)
hc <- hclust(d)
hc
```

Call:

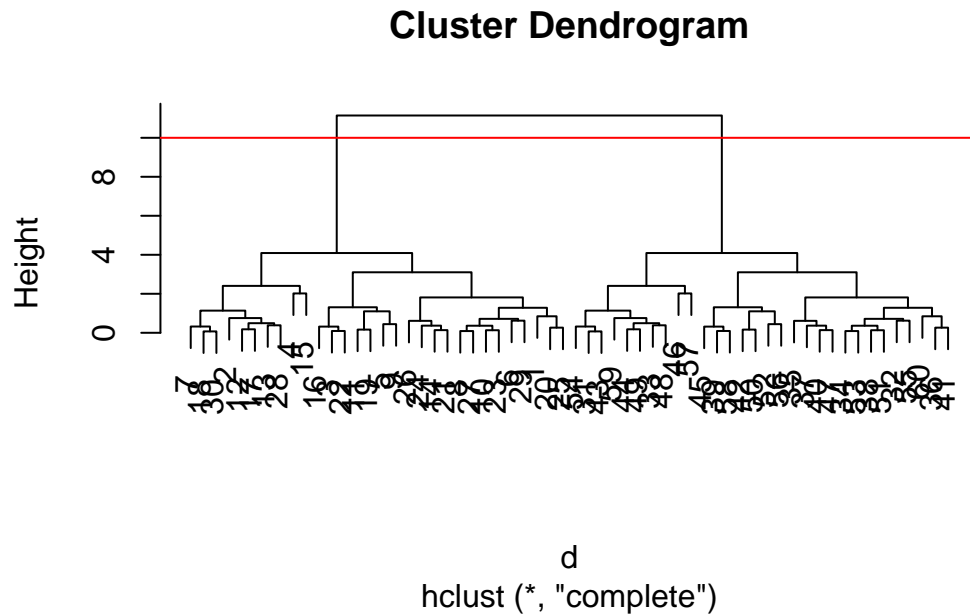
```
hclust(d = d)
```

```
Cluster method : complete
Distance       : euclidean
Number of objects: 60
```

Use `plot()` to view the results



```
plot(hc)
abline(h=10, col="red")
```




To make the “cut” and get our cluster membership vector we can use the `cutree()` function.

```
grps <- cutree(hc, h=10)
grps
```

```
[1] 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 2 2 2 2 2 2 2 2
[39] 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2
```

Make a plot of our data colored by hclust results

```
plot(x, col=grps)
```



## Principal Component Analysis (PCA)

Here we will do Principal Component Analysis (PCA) on some food data from the UK.

```
url <- "http://tinyurl.com/UK-foods"
x <- read.csv(url, row.names = 1)
View(x)
```

```
#rownames(x) <- x[,1]
#x <- x[,-1]
#x
```

```
#not a good method, overrides the functions
```

Q. How many rows and columns are in your new data frame named x? What R functions could you use to answer this question?

```
nrow(x)
```

```
[1] 17
```

```
ncol(x)
```

```
[1] 4
```

```
#or  
dim(x)
```

```
[1] 17 4
```

Preview the first six rows

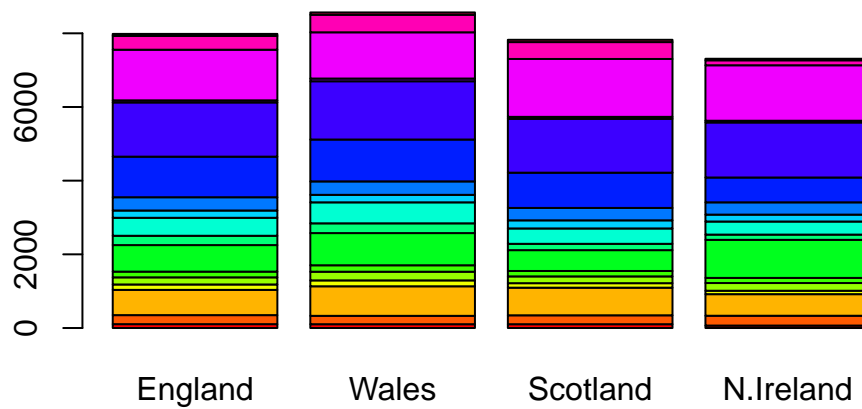
```
head(x)
```

	England	Wales	Scotland	N.Ireland
Cheese	105	103	103	66
Carcass_meat	245	227	242	267
Other_meat	685	803	750	586
Fish	147	160	122	93
Fats_and_oils	193	235	184	209
Sugars	156	175	147	139


Q2. Which approach to solving the **row-name** problems mentioned above do you prefer and why? Is one approach more robust than another under certain circumstances?

I would prefer to use the `dim()` or `nrow()` function and then adding `head()` as a second function, since the `x <- x[,-1]` function reloads and overrides functions so it is not as reliable as `dim()`.

```
barplot(as.matrix(x), beside=FALSE, col=rainbow(nrow(x)))
```



```
pairs(x, col=rainbow(10), pch=16)
```



##PCA to the rescue

The main “base” R function for PCA is called `prcomp()`.

```
pca <- prcomp( t(x))
summary(pca)
```

Importance of components:

	PC1	PC2	PC3	PC4
Standard deviation	324.1502	212.7478	73.87622	3.176e-14
Proportion of Variance	0.6744	0.2905	0.03503	0.000e+00
Cumulative Proportion	0.6744	0.9650	1.00000	1.000e+00

Q. How much variance is captured in 2 PCs?

96.5%

To make our main “PC score plot” (aka “PC1 vs PC2 plot” or “PC plot” or “ordination plot”).

```
attributes(pca)
```

\$names

```
[1] "sdev"      "rotation" "center"    "scale"     "x"
```

\$class


```
[1] "prcomp"
```

We are after the `pca$x` result component to make our main PCA plot.

```
pca$x
```

	PC1	PC2	PC3	PC4
England	-144.99315	-2.532999	105.768945	-4.894696e-14
Wales	-240.52915	-224.646925	-56.475555	5.700024e-13
Scotland	-91.86934	286.081786	-44.415495	-7.460785e-13
N.Ireland	477.39164	-58.901862	-4.877895	2.321303e-13

```
mycols <- c("orange", "red", "blue", "darkgreen")
plot(pca$x[,1], pca$x[,2], col=mycols, pch=16, xlab="PC1 (67.4%)", ylab="PC2 (29%)")
```



Another important result from PCA is how the original variables (in this case the foods) contribute to the PCs.

This is contained in the `pca$rotation` object - called the “loadings” or “contributions” to the PCs.

```
pca$rotation
```

	PC1	PC2	PC3	PC4
Cheese	-0.056955380	0.016012850	0.02394295	-0.694538519
Carcass_meat	0.047927628	0.013915823	0.06367111	0.489884628
Other_meat	-0.258916658	-0.015331138	-0.55384854	0.279023718
Fish	-0.084414983	-0.050754947	0.03906481	-0.008483145
Fats_and_oils	-0.005193623	-0.095388656	-0.12522257	0.076097502
Sugars	-0.037620983	-0.043021699	-0.03605745	0.034101334
Fresh_potatoes	0.401402060	-0.715017078	-0.20668248	-0.090972715
Fresh_Veg	-0.151849942	-0.144900268	0.21382237	-0.039901917
Other_Veg	-0.243593729	-0.225450923	-0.05332841	0.016719075
Processed_potatoes	-0.026886233	0.042850761	-0.07364902	0.030125166
Processed_Veg	-0.036488269	-0.045451802	0.05289191	-0.013969507
Fresh_fruit	-0.632640898	-0.177740743	0.40012865	0.184072217
Cereals	-0.047702858	-0.212599678	-0.35884921	0.191926714

Beverages	-0.026187756	-0.030560542	-0.04135860	0.004831876
Soft_drinks	0.232244140	0.555124311	-0.16942648	0.103508492
Alcoholic_drinks	-0.463968168	0.113536523	-0.49858320	-0.316290619
Confectionery	-0.029650201	0.005949921	-0.05232164	0.001847469

Now to make a plot around PC1

```
library(ggplot2)

contrib <- as.data.frame(pca$rotation)

ggplot(contrib) +
  aes(PC1, rownames(contrib))+ geom_col(fill="blue")
```

